

Multilayer Perceptron Architecture

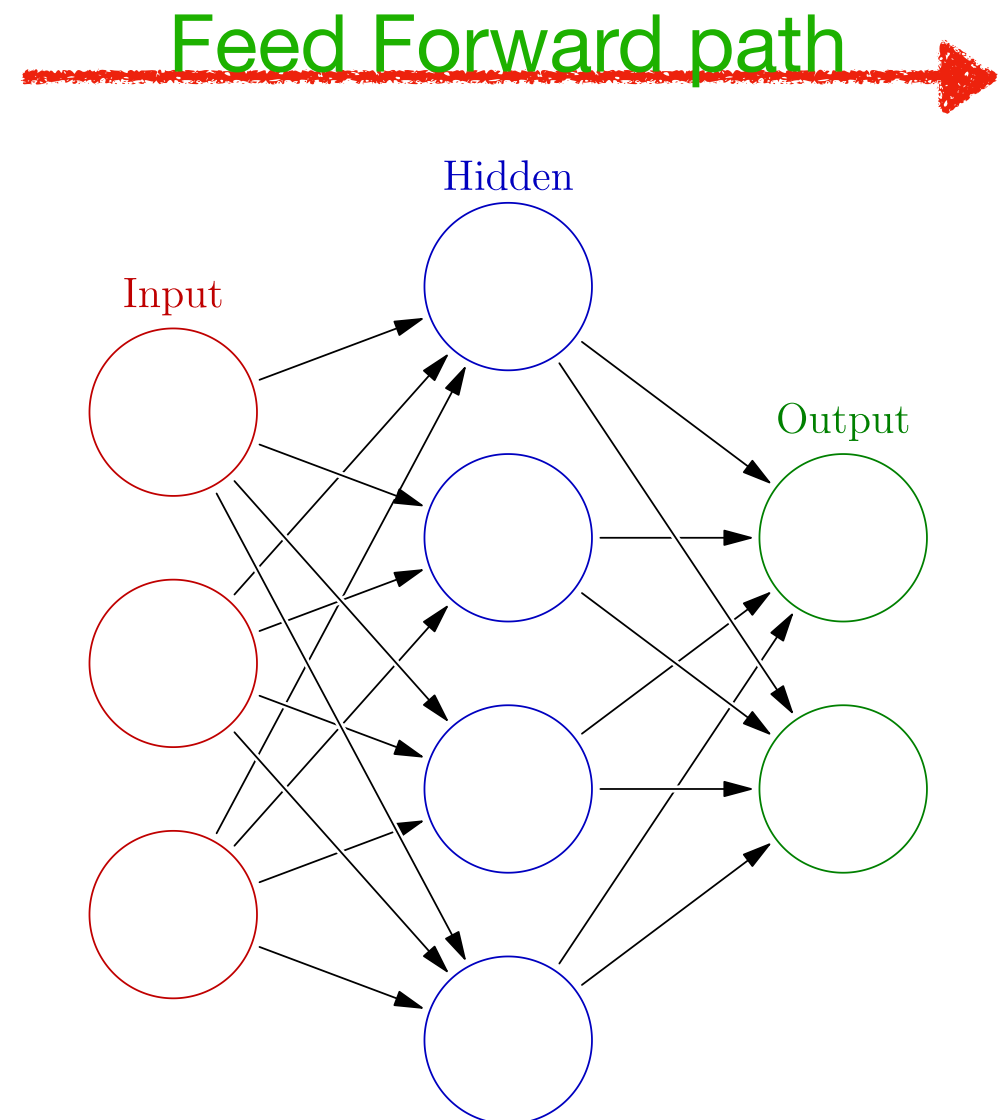
Understanding Neural Network Layers

Outline

- Feedforward Network Basics
- Input Layer Fundamentals
- Hidden Layers Exploration
- Output Layer Design
- Network Complexity Analysis
- Mathematical Representation
- Architectural Considerations
- Practical Implementation


Feedforward Network Basics

- Information flows from input to the output (unidirectional)
- No Loops/Cycles
- No Memory of Past Inputs (unlike RNNs)




Input Layer

- Main Purpose:
 - Take in the raw data
 - Ensures input shape is compatible with the architecture
- Doesn't transform the features
- Each feature/dimension of the input data corresponds to a node in input layer
- It distributes the input data to all the nodes in the first hidden layer (each input node is linked to every node in the first hidden layer)




Hidden Layer

- Lies between the input and the output layer
- Feature Extraction
 - The first hidden layer captures the basic features (e.g. edges in an image)
 - Deeper layers, extract more complex, abstract features (e.g. shapes, eyes, faces)
 - Eliminates the need for manual feature engineering
- # hidden layers = depth of the network
- # neurons/nodes per layer = Width of the network
- Networks with many hidden layers = Deep neural networks (DNNs)




Hidden Layer

- Tasks requiring abstract reasoning
 - **Depth** is important
- Tasks requiring fine-grained feature analysis
 - **Width** is important




Output Layer

- Final Layer of the network
- Produces the network predictions
- Structure
 - **Regression**
 - **Scalar Output:** Single Node
 - **Vector Output:** Multiple Nodes
 - **Classification**
 - # Nodes = # Classes
- Output values must often be post-processed to get a meaningful result






Output Layer

- Some Designs:
 - **Single Node** (Regression)
 - **Single Node** (Binary Classification)
 - **Multi-Node** (Multi-Class Classification)
 - **Multi-Node** (Multi-Label Classification)
 - **Custom**




Overfitting vs. Underfitting




Complexity Trade-Offs

- Between the ability to learn and practical limits (i.e. Computation, Generalization, Interpretability)
- **Complexity vs. Generalization**
 - **High Complexity:** Learns well, Risk of overfitting
 - **Low Complexity:** Learns poorly, Risk of underfitting
- **Depth vs. Width**
 - **Deeper Networks:** Learn complex features, Risk of Vanishing/Exploding Gradients
 - **Wider Network:** Excel in capturing fine-grained patterns, But can be computationally expensive, Risk of Overfitting




Complexity Trade-Offs

- Between the ability to learn and practical limits (i.e. Computation, Generalization, Interpretability)
- **Accuracy vs. Efficiency**
 - **High Accuracy:** Requires larger models, Increase computation, memory usage and energy consumption
 - **Efficiency:** Cut some part of the model, lose some accuracy in favor of faster inference and reduce resource requirements
- **Interpretability vs. Complexity**
 - **Complex model:** Nice but like a Black box
 - **Simple model:** Weak in complex tasks but interpretable




Complexity Trade-Offs

- Between the ability to learn and practical limits (i.e. Computation, Generalization, Interpretability)
- **Robustness vs. Simplicity**
 - **Robust Model:** Account for noisy or adversarial data, but computationally intensive and complex
 - **Simple Model:** Struggles with Noisy data but faster and easier to implement




Strategies in Managing the Trade-Offs

- **Regularization**
- **Model Pruning**
- **Early Stopping**
- **Transfer Learning**
- **Automated Optimization**
 - **Neural Architecture Search (NAS)**



Mathematics of MLP

- Each layer has neurons, with every neuron in that layer, linked to the neurons of the next
- For a given layer *l*
 - Input: x^{l-1}
 - Linear Transformation:
$$z^l = \omega^{(l)} x^{(l-1)} + b^{(l)}$$
 - Non-Linear Transformation:
$$x^{(l)} = \phi(z^{(l)})$$




Mathematics of MLP

- **Weights:**

- A matrix of learnable parameters $\omega^{(l)} \in \mathbb{R}^{n_l \times n_{l-1}}$
- Represents the strength of connections between neurons in layer “ $l-1$ ” and “ l ”

- **Biases:**

- A vector of learnable parameters, where $b^{(l)} \in \mathbb{R}^{n_l}$
- Allows the layer to shift the activation values → Enhancing flexibility



Mathematics of MLP


- **Activations:**

- **ReLU:** $\phi(x) = \max(0, x)$

- **Sigmoid:** $\phi(x) = \frac{1}{1 + e^{-x}}$

- **Tanh:** $\phi(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$

- **Softmax:** $\phi(x_i) = \frac{e^{x_i}}{\sum_j e^{x_j}}$



Mathematics of MLP

- **Forward Propagation:**

- ***Input Layer***


- $x^{(0)} = input$

- ***Hidden Layer:***

- $x^l = \phi(\omega^{(l)}x^{(l-1)} + b^l)$


- ***Output Layer***

- $\hat{y} = \phi(\omega^{(L)}x^{(L-1)} + b^L)$




Considerations

- **Input Layer**
 - **Size:** ~# features in input data
- **Hidden Layer**
 - *Determine the network's depth and width*
 - **Shallow Networks:** Suitable for simpler tasks or smaller datasets.
 - **Deep Networks:** Capture hierarchical relationships but require careful regularization and optimization.
- **Output Layer:**
 - **Size:** # output classes
 - *Activation functions should align with the task type (e.g., softmax for classification, linear for regression).*




Considerations

- **Neurons per Layer**
 - **Too Few Neurons:** Can lead to underfitting, due to lack of ability in learning complex patterns
 - **Too Many Neurons:** Can lead to overfitting, model memorizes the training data instead of generalizing
 - **Practice:** For a funnel-shaped structure with fewer neurons in deeper layers
- **Regularization**
 - **Dropout**
 - **Weight Decay (L2 Regularization)**
 - **Batch Normalization**
 - **Early Stopping**
 - **Data Augmentation**



Practice



```
class MLP(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super(MLP, self).__init__()
        self.layer1 = nn.Linear(input_size, hidden_size)
        self.layer2 = nn.Linear(hidden_size, output_size)
        self.activation = nn.ReLU()
        self.output_activation = nn.Softmax(dim=1) # For multi-class probabilities

    def forward(self, x):
        x = self.activation(self.layer1(x))
        x = self.output_activation(self.layer2(x))
        return x
```