

Learning Rates and Weight Initialization


Foundations of Effective Neural Network Training

Outline

- Learning Rate Fundamentals
- Learning Rate Impact
- Weight Initialization Basics
- Common Initialization Strategies
- Vanishing Gradient Problem
- Exploding Gradient Problem
- Initialization Techniques

Learning Rate Fundamentals


- **Learning Rate (η)** is a critical hyperparameter in optimization
 - It matters how you approach the local minima
- **Small Learning Rate**
 - Slow Convergence
 - Can get stuck in minimas like saddle points
- **Large Learning Rate**
 - Fast Convergence, but might overshoot the desired minimum
 - Can cause Loss function to oscillate/diverge
- **Optimal Learning Rate**
 - Balances the speed and stability



Learning Rate Fundamentals


- **Challenges**

- Choosing the right value
- Dynamic Adjustments
 - Learning Rate schedules(i.e, step decay, exponential decay)
- Sensitivity to scale
 - Different features or layers may need different learning rates
 - Feature Scaling
- Impact of loss surface
 - A highly non-convex loss surface with multiple local minima and saddle points complicates the choice of learning rate.




Weight Initialization Basics

- Should we initialize the weights randomly?!




Weight Initialization Basics

- Process of assigning initial values to weights
 - It matters where we start from
- Proper Initialization is crucial
 - Ensure Effective Signal Propagation
 - Improve convergence speed
 - Prevent symmetry




Weight Initialization Basics

- When **Randomly** initialize the weights
- There is a chance to have bad start points
- **Vanishing Gradients**
- **Exploding Gradients**
- **Symmetry Breaking**
- **Unstable Learning Dynamics**



Impact of Weight initialization

- Training Speed
- Model Accuracy
- Prevention of Overfitting



Common Initialization Strategies

- **Random Initialization**

$$\omega \sim \text{Uniform}(-\epsilon, \epsilon)$$

- **Xavier/Glorot Initialization**


$$\omega \sim \text{Uniform}\left(-\sqrt{\frac{1}{n}}, \sqrt{\frac{1}{n}}\right)$$

- **He Initialization**

$$\omega \sim \mathcal{N}\left(0, \sqrt{\frac{2}{n}}\right) \text{ or } \omega \sim U\left(-\sqrt{\frac{6}{n}}, \sqrt{\frac{6}{n}}\right)$$

- **Zero Initialization**

- Should be avoided for weights but can be used for bias initialization



Vanishing Gradient Problem

- It happens when gradients become extremely small during backpropagation
- **Root Causes**
 - **Activation Functions**
 - Saturating functions like **sigmoid** and **Tanh**
 - In case of Sigmoid: $f'(x) = f(x)(1 - f(x))$, becomes very small if $f(x)$ is close to 0 or 1
 - **Weight Initialization**
 - **Deep Networks**
 - **Loss Function**
 - Functions with **small derivatives** like the ones used in **classification tasks**

Vanishing Gradient Problem

- **Impacts on Deep Network**
 - Slow Learning in early layers
 - Poor model performance
 - Optimization difficulties
 - Bias toward output layers

How To Identify Vanishing Gradient

- **Gradient Magnitude Inspection**
 - Measure the magnitude of gradients at different layers during training
 - The early layers will have near-zero gradient values
- **Training Behavior**
 - Slow Convergence
- **Activation Distribution**
 - Saturation near min/max of the activation range, suggests potential gradient vanishing
- **Weight Updates**
 - Small changes in weights of earlier layers during training

How To Mitigate Vanishing Gradient

- **Use NonSaturating Functions like ReLU**
- **Weight Initialization Techniques**
- **Batch Normalization**
- **Residual Networks?!**

Exploding Gradient Problem

- It happens when gradients grow uncontrollably large during backpropagation
- **Root Causes**
 - **Deep Networks**
 - If gradients in different layers are larger than 1, then the repeated multiplication of them would be huge
 - **Recurrent Neural Networks (RNNs)**
 - Gradients are propagated through time and long sequences amplify the effect
 - **Improper weight Initialization**
 - **Activation Functions with unbounded outputs**

How To Identify Exploding Gradient

- **Gradient Magnitude Inspection**
 - Measure the magnitude of gradients at different layers during training
- **Training Behavior**
 - Observe instabilities, Diverging loss values, NaN values
- **Activation Distribution**
- **Log Gradient Values**
 - Visualize gradients using **TensorBoard**. Even a sudden spike is a sign.

How To Mitigate Vanishing Gradient

- **Gradient Clipping**

$$\text{If } \|\nabla L\| > \tau, \nabla L \leftarrow \tau \cdot \frac{\nabla L}{\|\nabla L\|}$$

- **Weight Regularization**

- **L2** regularization penalizes the large weights

- **Proper Weight Initialization**

- **Xavier/Glorot** or **He** initialization ensure the weights are scaled properly

- **Optimizers with adaptive learning**

- **Adam** or **RMSProp**, dynamically adjust the learning rate for each parameter

- **Batch Normalization**

- **Architectural Adjustments**

- For **RNNs**, you might use **LSTMs** or **GRUs** which include gating mechanisms to control the gradient flow

Initialization Techniques

- **Zero Initialization**
 - **Pros:** Simple, and works for biases
 - **Cons:** All neurons in a layer learn the same features because their gradients are identical
- **Random Initialization**
 - **Pros:** Breaking Symmetry and Enables effective learning
 - **Cons:** Randomly chosen values can result in vanishing/exploding gradients
- **He Initialization**
 - **Pros:** Prevents exploding/vanishing gradients
 - **Cons:** Might not be optimal for activations **other than ReLU**, and a slightly more computation overhead

Initialization Techniques

Method	Pros	Cons	Best for
Zero Initialization	Simple	Symmetry	Bias Initialization
Random Initialization	Breaks Symmetry	May cause vanishing/ exploding gradients	Shallow networks
He Initialization	Prevents gradient issues	Slight computational overhead	Deep networks with ReLU