

Performance Metrics

What is Good Performance?

Outline

- Classification Metrics
- Regression Metrics
- Segmentation Metrics
- Computational Efficiency Metrics
- Training and Optimization Metrics

Classification Metrics

- **Accuracy**

- The ratio of correctly predicted labels to the total number of predictions

$$Accuracy = \frac{\# \text{ correct prediction}}{\# \text{ Total Predictions}}$$

- Suitable for **balanced datasets** but less informative for imbalanced datasets

- **Precision**

- Measures the proportion of true positive predictions out of all positive predictions

$$Precision = \frac{True\ Positive(TP)}{True\ Positive(TP) + False\ Positive(FP)}$$

- Important for imbalanced datasets where false positives or false negatives carry significant weight

- **Recall (sensitivity)**

- Measures the proportion of true positives correctly identified out of all actual positives

$$Recall = \frac{True\ Positive(TP)}{True\ Positive(TP) + False\ Negative(FN)}$$

- Important for imbalanced datasets where false positives or false negatives carry significant weight

Classification Metrics

- **F1-Score**

- The harmonic mean of precision and recall

$$F1 - Score = 2 \cdot \frac{Precision \cdot Recall}{Precision + Recall}$$

- Important for imbalanced datasets where false positives or false negatives carry significant weight

- **ROC-AUC (Receiver Operating Characteristic - Area Under Curve)**

- Evaluates the trade-off between true positive rate (TPR) and false positive rate (FPR) across different thresholds
- Useful for **binary classification** and evaluating **model discrimination capability**

Regression Metrics

- **Mean Squared Error (MSE)**

- Measures the average squared difference between predicted and actual values

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

- Penalizes larger errors more than smaller ones

- **Mean Absolute Error (MAE)**

- Measures the average absolute difference between predicted and actual values

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

- Provides a more interpretable measure compared to MSE

- **R^2 (Coefficient of Determination)**

- Measures how well the predictions explain the variance in the actual data

$$R^2 = 1 - \frac{\sum (y_i - \hat{y}_i)^2}{\sum (y_i - \bar{y})^2}$$

- Indicates the goodness of fit for regression tasks

Segmentation Metrics

- **Intersection over Union (IoU)**

- Measures the overlap between predicted and ground truth regions

$$IoU = \frac{\textit{Area of Overlap}}{\textit{Area of Union}}$$

- Common in image segmentation tasks to evaluate boundary precision

- **Dice Coefficient**

- Measures the similarity between predicted and ground truth masks

$$Dice = \frac{2 \cdot |predicted \cap Ground Truth|}{|predicted| + |Ground Truth|}$$

- Particularly useful for medical image analysis

- **Pixel Accuracy**

- Proportion of correctly classified pixels

$$Pixel Accuracy = \frac{\textit{Correct Pixels}}{\textit{Total Pixels}}$$

Computational Efficiency Metrics

- **FLOPs (Floating Point Operations Per Second)**
 - Measures the computational complexity of the model by counting the number of operations
- **Inference Time**
 - Time taken for a model to process input data and produce an output
- **Model Size**
 - Measures the storage requirement for the model, typically in megabytes (MB) or gigabytes (GB)
 - Or in terms of **# of Parameters**

Training and Optimization Metrics

- **Loss Function**

- Tracks the error during training
 - **Cross-Entropy Loss:** For classification tasks
 - **Mean Squared Error:** For regression tasks

- **Learning Curves**

- Plots of training and validation loss/accuracy over epochs, used to monitor overfitting or underfitting